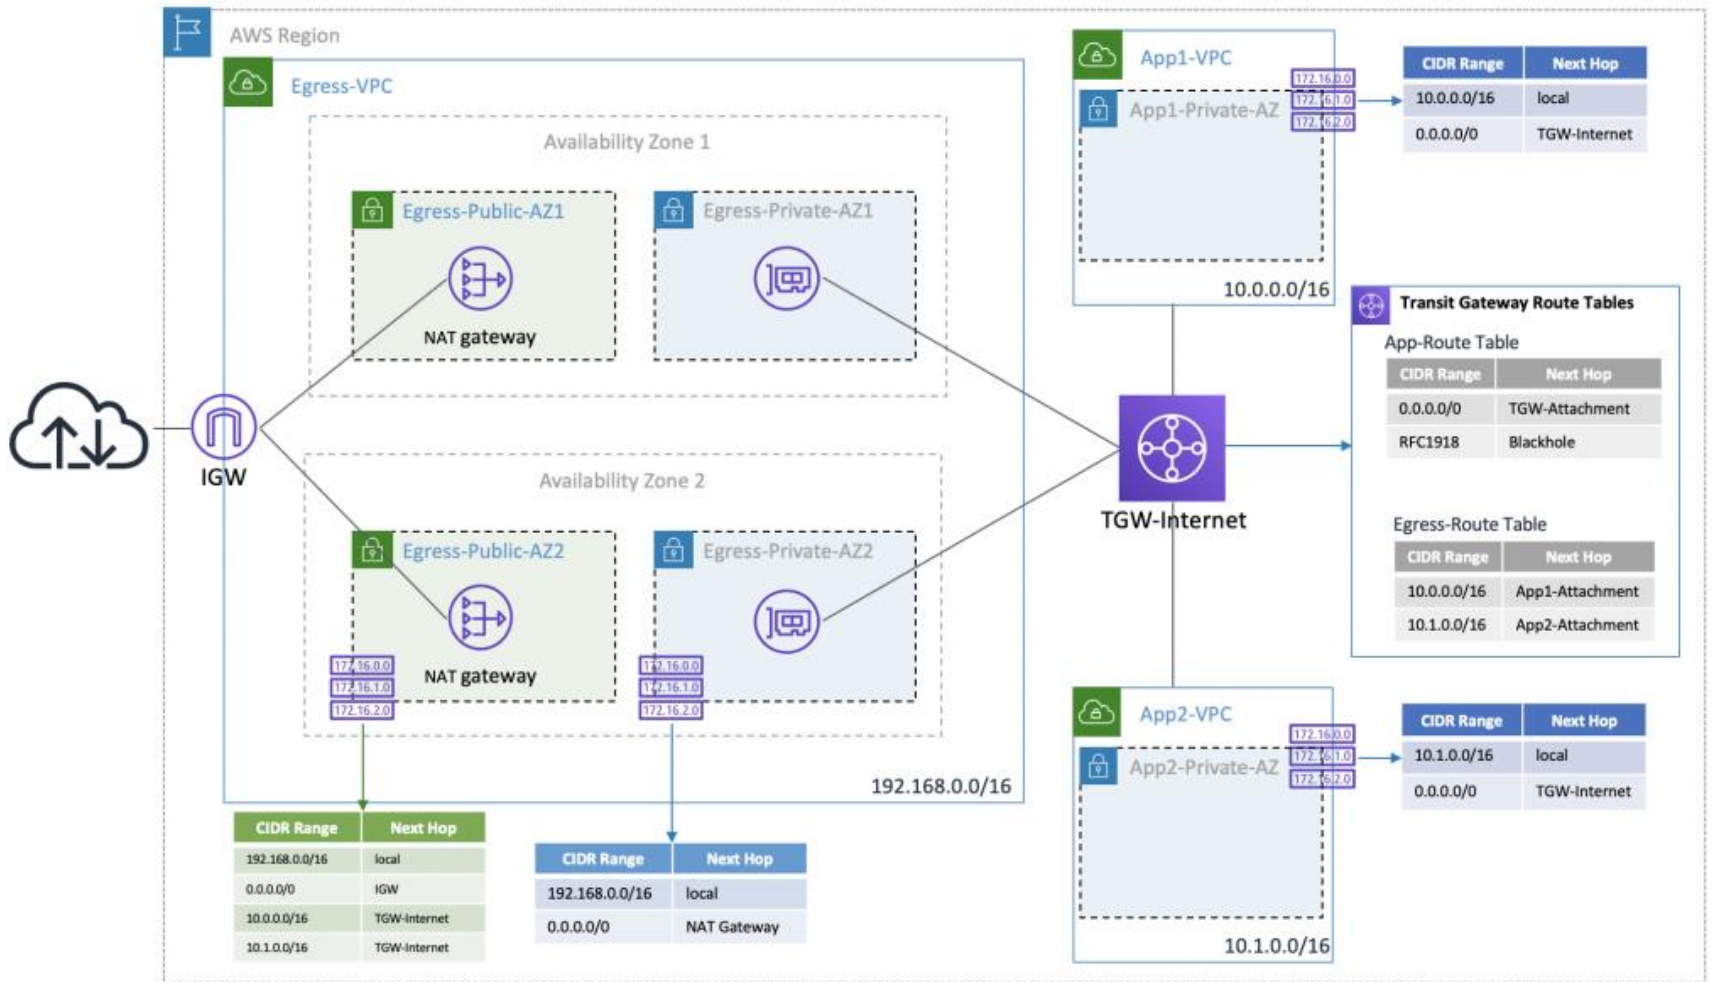


CSAA Practice Test-4

Interface Endpoint

Gateway Endpoint


What	<u>Elastic Network Interface</u> with a Private IP	A gateway that is a target for a specific <u>route</u>
How	Uses DNS entries to redirect traffic	Uses prefix lists in the route table to redirect traffic
Which services	API Gateway, CloudFormation, CloudWatch etc.	<u>Amazon S3, DynamoDB</u>
Security	Security Groups	VPC Endpoint Policies



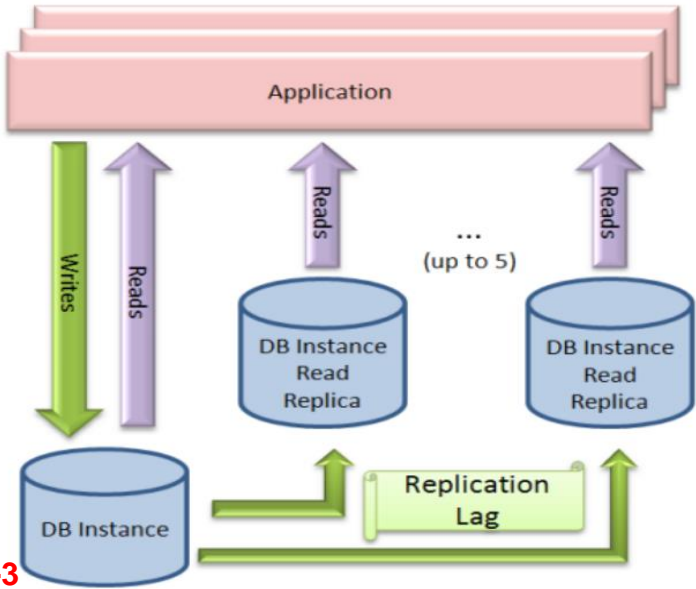

<https://aws.amazon.com/blogs/networking-and-content-delivery/creating-a-single-internet-exit-point-from-multiple-vpcs-using-aws-transit-gateway>

You can have Multi-AZ replicas, replicas in another region, or even replicas of other read replicas


Multi-AZ Replicas



Cross-Region Replicas



- When bad queries (such as lack of primary or unique keys) are replicated
- Some trouble with the hardware (such as network or disk IO issues).



Create role

▼ Attach permissions policies

Choose one or more policies to attach to your new role.

Create policy

Filter policies ▼

Q SQS

Policy name ▼

- ☐ ► AmazonSQSFullAccess
- ☐ ► AmazonSQSReadOnlyAccess
- ☐ ► AWSLambdaSQSQueueExecutionRole

Choose trail attributes

General details

A trail created in the console is a multi-region trail. [Learn more](#)

Trail name

Enter a display name for your trail.

3-128 characters. Only letters, numbers, periods, underscores, and dashes are allowed.

☐ Enable for all accounts in my organization

To review accounts in your organization, open AWS Organizations. [See all accounts](#)

Storage location [Info](#)



Create new S3 bucket

Create a bucket to store logs for the trail.



Use existing S3 bucket

Choose an existing bucket to store logs for this trail.

Trail log bucket and folder

Enter a new S3 bucket name and folder (prefix) to store your logs. Bucket names must be globally unique.

Logs will be stored in aws-cloudtrail-logs-827784331229-59a5cdbe/AWSLogs/827784331229

Log file SSE-KMS encryption [Info](#)

☒ Enabled

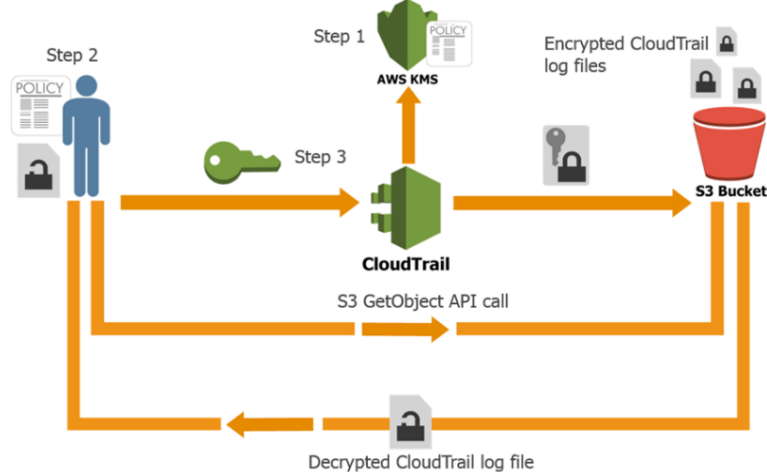
AWS KMS customer managed CMK

☒ New

☐ Existing

AWS KMS alias

KMS key and S3 bucket must be in the same region.



S3 Data Consistency

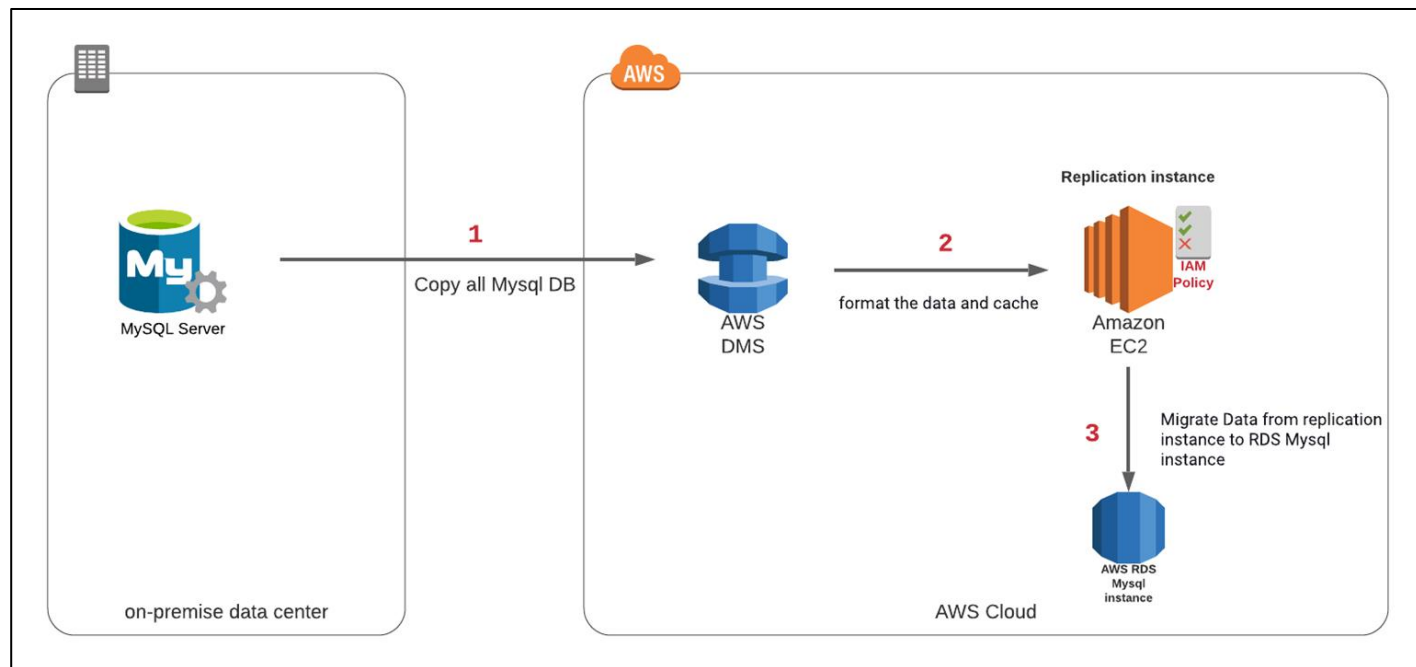
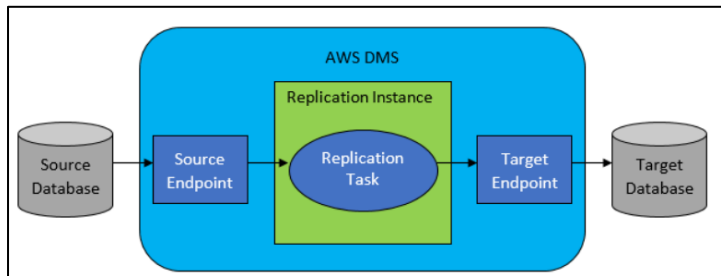
Strong read-after-write consistency

- PUT and DELETE requests of objects
- read operations on Amazon S3 Select, Amazon S3 access controls lists (ACLs), Amazon S3 Object Tags, and object metadata (for example, the HEAD object)

Eventual consistency

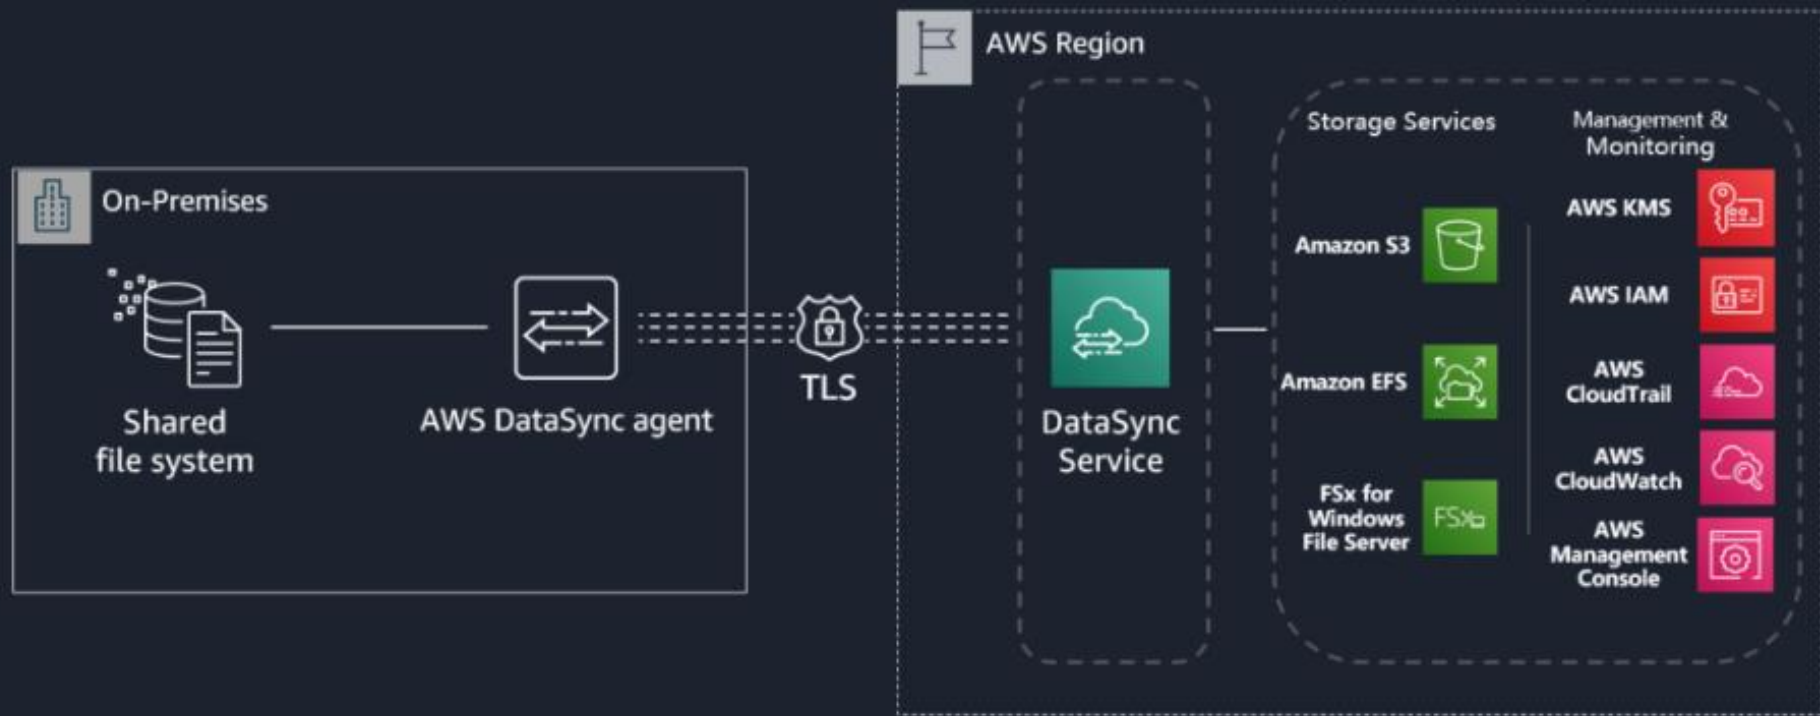
- Bucket configurations
- can be a slightly delay

- Cross-origin resource sharing (CORS) defines a way for client web applications that are loaded **in one domain to interact with resources in a different domain.**
- In certain cases, the developer of the original page might have legitimate reasons to write code **that interacts with content or services at other locations.** CORS provides the mechanism to allow the developer to tell the browser to allow this interaction.
- To configure your bucket to allow cross-origin requests, you create a CORS configuration. The CORS configuration is a document with rules that identify the origins that you will allow to access your bucket, the operations (HTTP methods) that will support for each origin, and other operation-specific information.



How does AWS DataSync work?

Simplifies, automates, and accelerates data transfer to or from AWS



Create Distribution

Origin Settings

Origin Domain Name


Origin Path

Enable Origin Shield ☐ Yes
☒ No

Origin ID

Restrict Bucket Access ☒ Yes
☐ No

Origin Access Identity ☒ Create a New Identity
☐ Use an Existing Identity



signed URLs or
signed cookies
(https only)

1

i

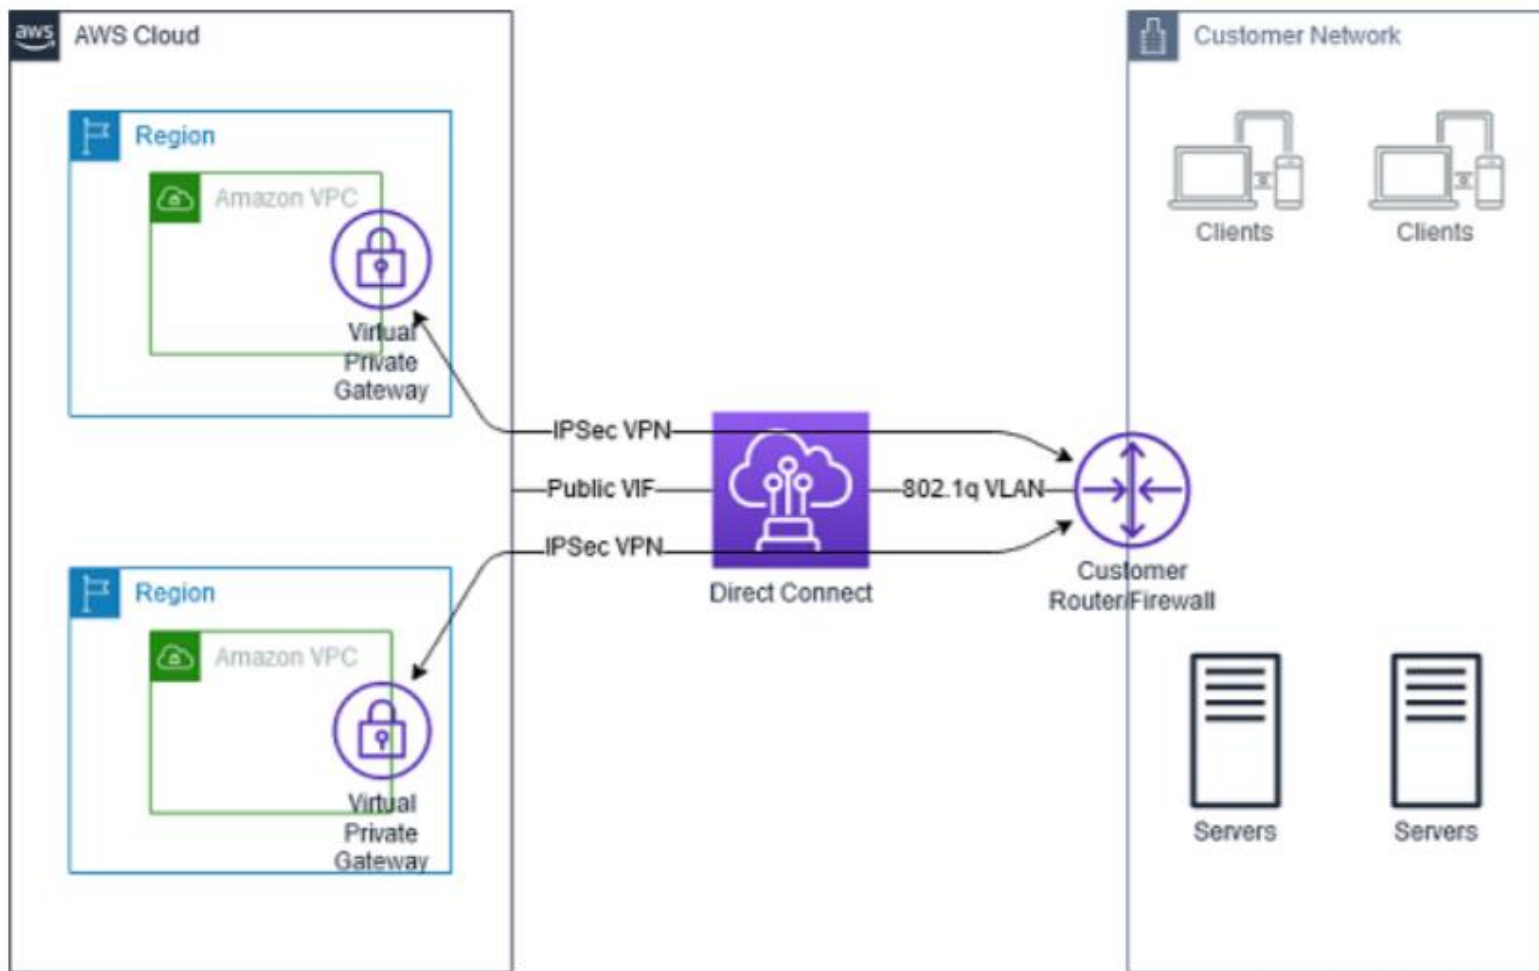
i

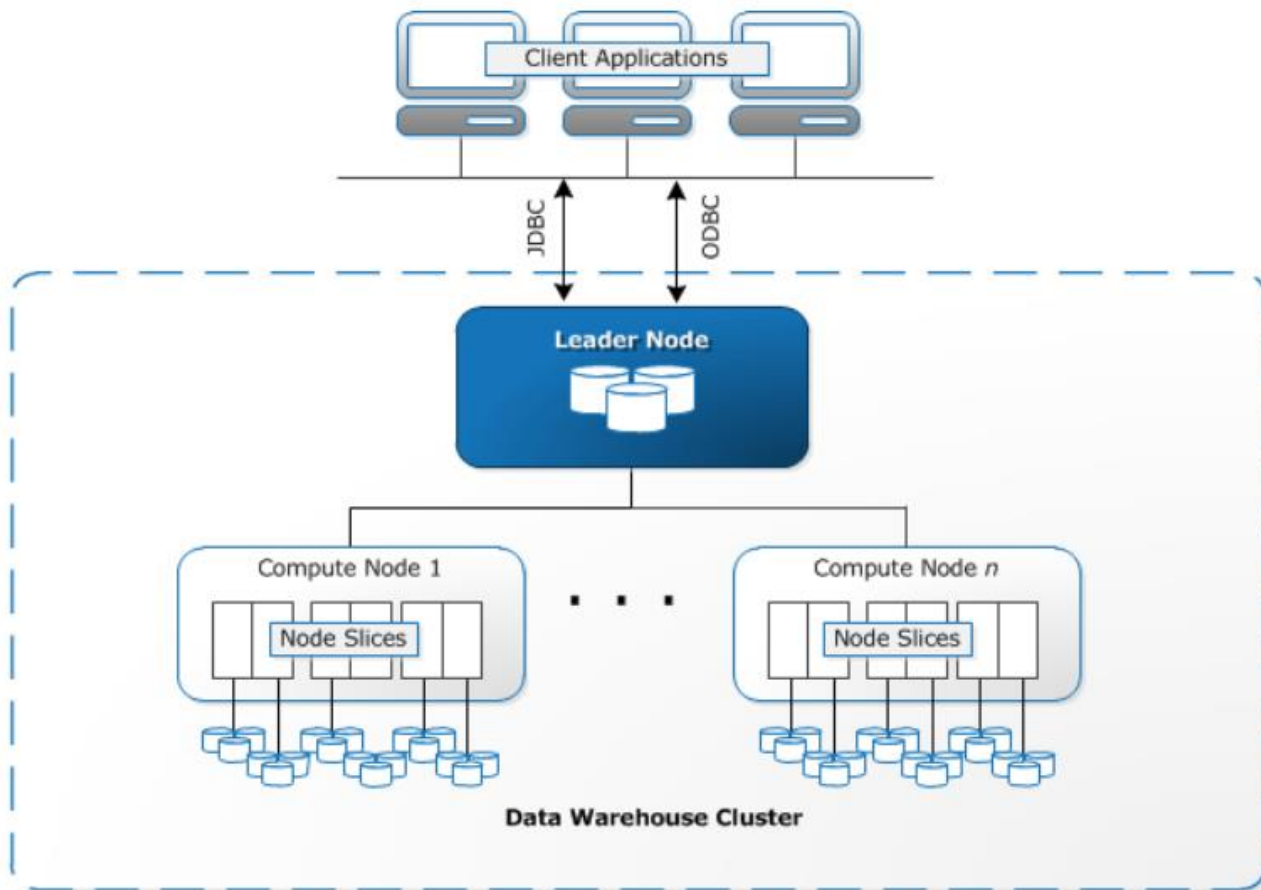
Enter a description for the origin. This value lets you distinguish multiple origins in the same distribution from one another. The description for each origin must be unique within the distribution.

i

If you want to require that users always access your Amazon S3 content using CloudFront URLs, not Amazon S3 URLs, click Yes. This is useful when you are using signed URLs or signed cookies to restrict access to your content. In the Help, see "Serving Private Content through CloudFront".

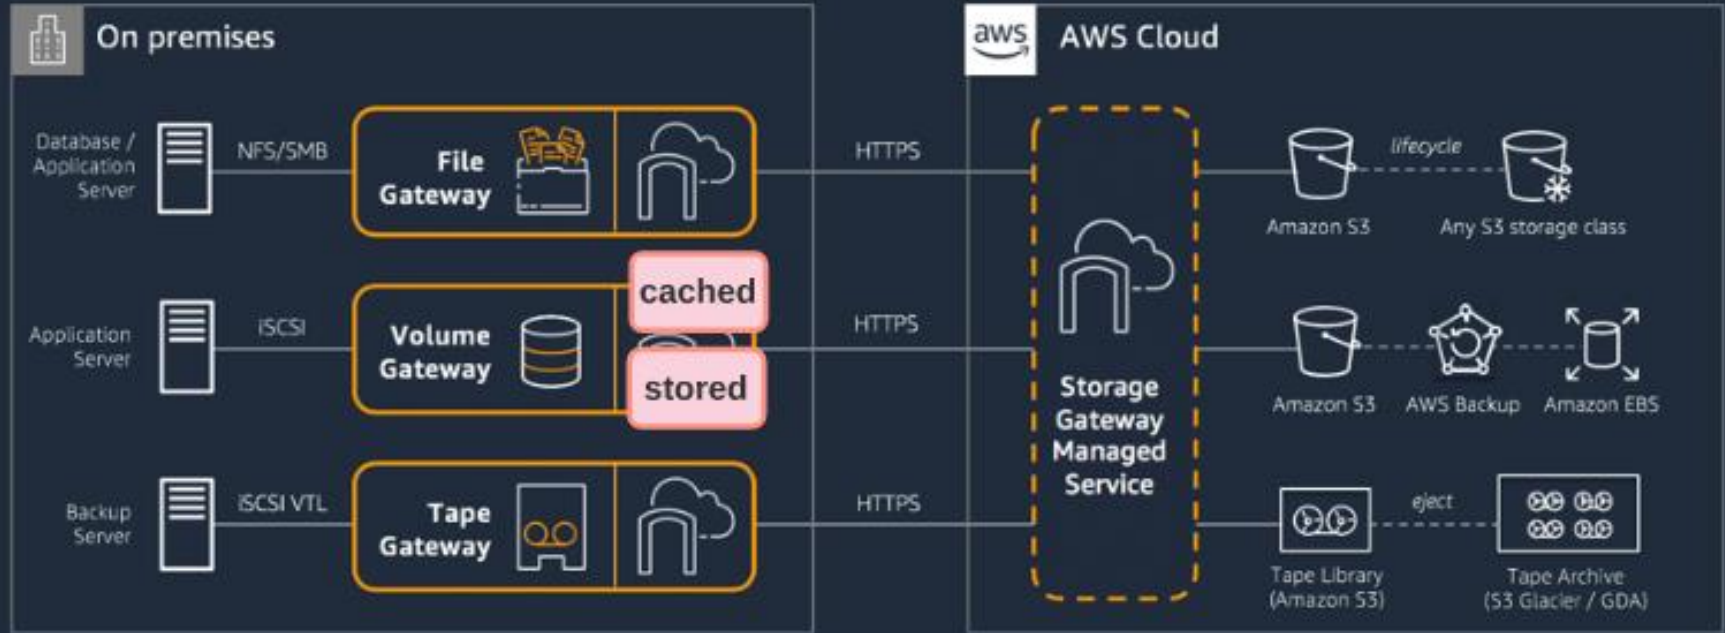
i





Move on-premises backups to the cloud

Maintain your backup workflows while reducing your backup infrastructure on-premises



Data on the volumes is stored in Amazon S3 and you can take point in time copies of volumes which are stored in AWS as Amazon EBS snapshots.

lambda > Functions > xxx

XXX

ConfigurationPermissionsMonitoring

▼ Designer

+ Add trigger

xxx

Layers (0)

Trigger configuration

Select a trigger

API Gateway
api application-services aws serverless

AWS IoT
aws devices iot

Alexa Skills Kit
alexa iot

Alexa Smart Home
alexa iot

Application Load Balancer
aws load-balancing

CloudFront
aws cdn edge

CloudWatch Logs
aws logging management-tools

CodeCommit
aws developer-tools git

Cognito Sync Trigger
authentication aws identity mobile-services sync

DynamoDB
aws database nosql

EventBridge (CloudWatch Events)
aws events management-tools

Kinesis
analytics aws streaming

MQ
aws messaging multi-protocol

MSK

Add trigger

Trigger configuration

API Gateway
api application-services aws serverless

Add an API to your Lambda function to create an HTTP endpoint that invokes your function. API Gateway supports two types of RESTful APIs: HTTP APIs and REST APIs. [Learn more](#)

API
Create a new API or attach an existing one.
Create an API

API type
☒ HTTP API
Create an HTTP API.
☐ REST API
Create a REST API.

Security
Configure the security mechanism for your API endpoint.

▼ Additional settings

API name
Choose a name for your API. API names don't need to be unique.
xxx-API

Deployment stage
The name of your API's deployment stage.
default

☒ Cross-origin resource sharing (CORS)
CORS is required to call your API from a webpage that isn't hosted on the same domain. This option enables cross-origin resource sharing (CORS) from any domain by adding the Access-Control-Allow-Origin header to all responses.

The AWS modern application platform

SERVERLESS MICROSERVICES



SERVERLESS DATASTORES



DEVELOPER TOOLS



SECURITY AND COMPLIANCE



Glacier Retrievals: Expedited and Bulk Retrievals

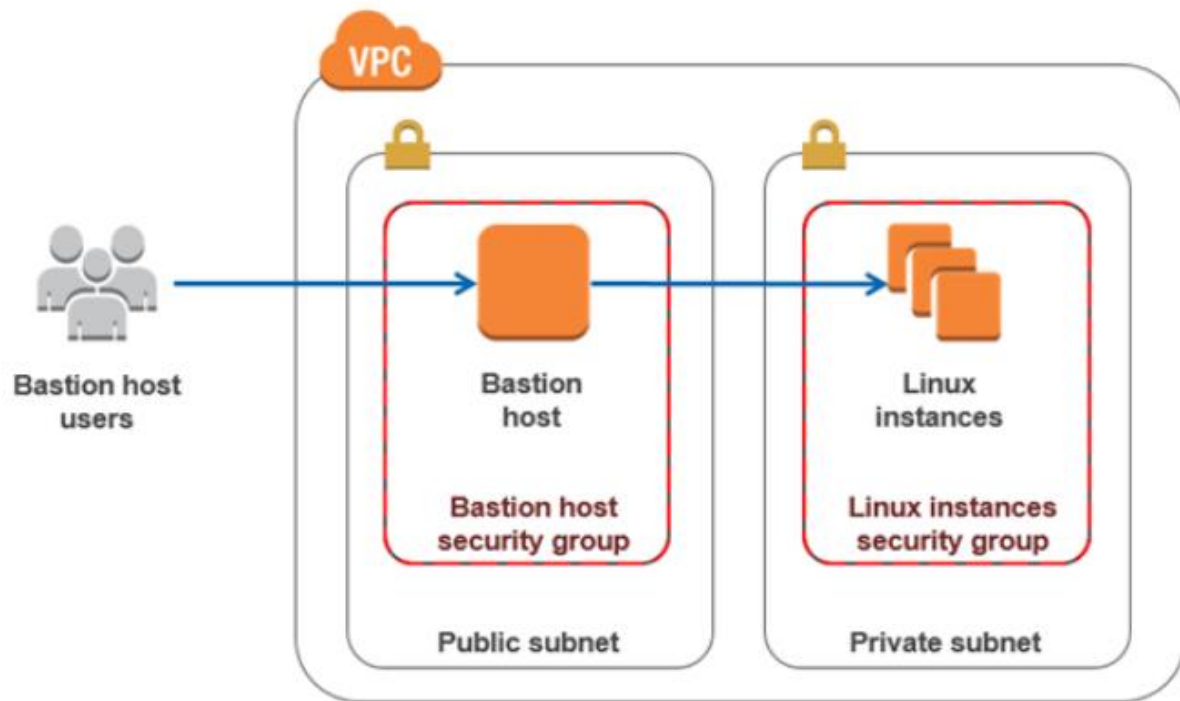
New

Clip slide

- Expedited: designed for occasional urgent access to a small number of archives
- Standard: Low-cost option for retrieving data in just a few hours
- Bulk: Lowest cost option optimized for large retrievals, up to petabytes of data in 12 hours
- Three flexible and powerful retrieval options to access any of your Glacier data

	Expedited	Standard	Bulk
Data Access Time	1 - 5 minutes	3 - 5 hours	5 - 12 hours
Data Retrievals	\$0.03 per GB	\$0.01 per GB	\$0.0025 per GB
Retrieval Requests	\$0.01 per request	\$0.05 per 1,000 requests	\$0.025 per 1,000 requests

aws



Types of Savings Plans



EC2 Instance Savings Plans

Provide the deepest discounts, up to 72% (same as Standard RIs) on the selected instance family (e.g. C5 or M5), in a specific AWS region

FLEXIBLE ACROSS

- ✓ Size: E.g. move from m5.xl to m5.4xl
- ✓ OS: E.g. change from m5.xl Windows to m5.xl Linux
- ✓ Tenancy: E.g. modify m5.xl Dedicated to m5.xl Default tenancy



Compute Savings Plans

Offer the greatest flexibility, up to 66% discounts (same discounts as Convertible RIs)

FLEXIBLE ACROSS

- ✓ Instance family: E.g. Move from C5 to M5
- ✓ Region: E.g. change from EU (Ireland) to EU (London)
- ✓ OS: E.g. Windows to Linux
- ✓ Tenancy: E.g. switch Dedicated tenancy to Default tenancy
- ✓ Compute options: E.g. move from EC2 to Fargate

Summary


Routes

Subnet Associations

Route Propagation

Tags

Edit

 Save Successful

View:

All rules ▼

Destination	Target	Status	Propagated
10.0.0.0/16	local	Active	No
0.0.0.0/0	igw-a97272cc	Active	No

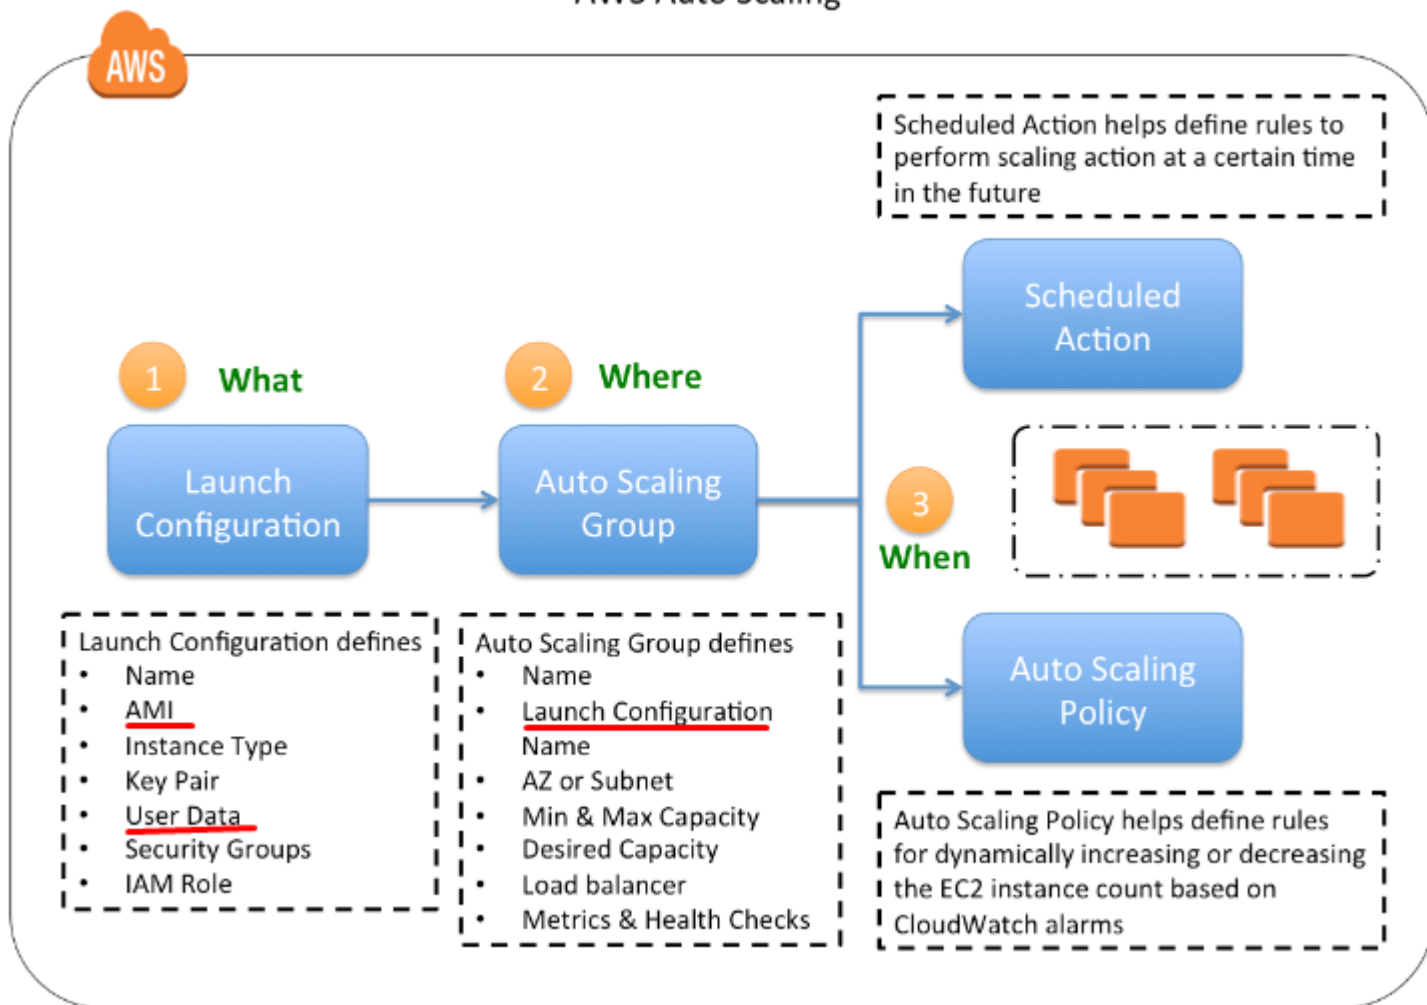
The destination for the route is 0.0.0.0/0, which represents **all IPv4 addresses**. The target is the internet gateway that's attached to your VPC.

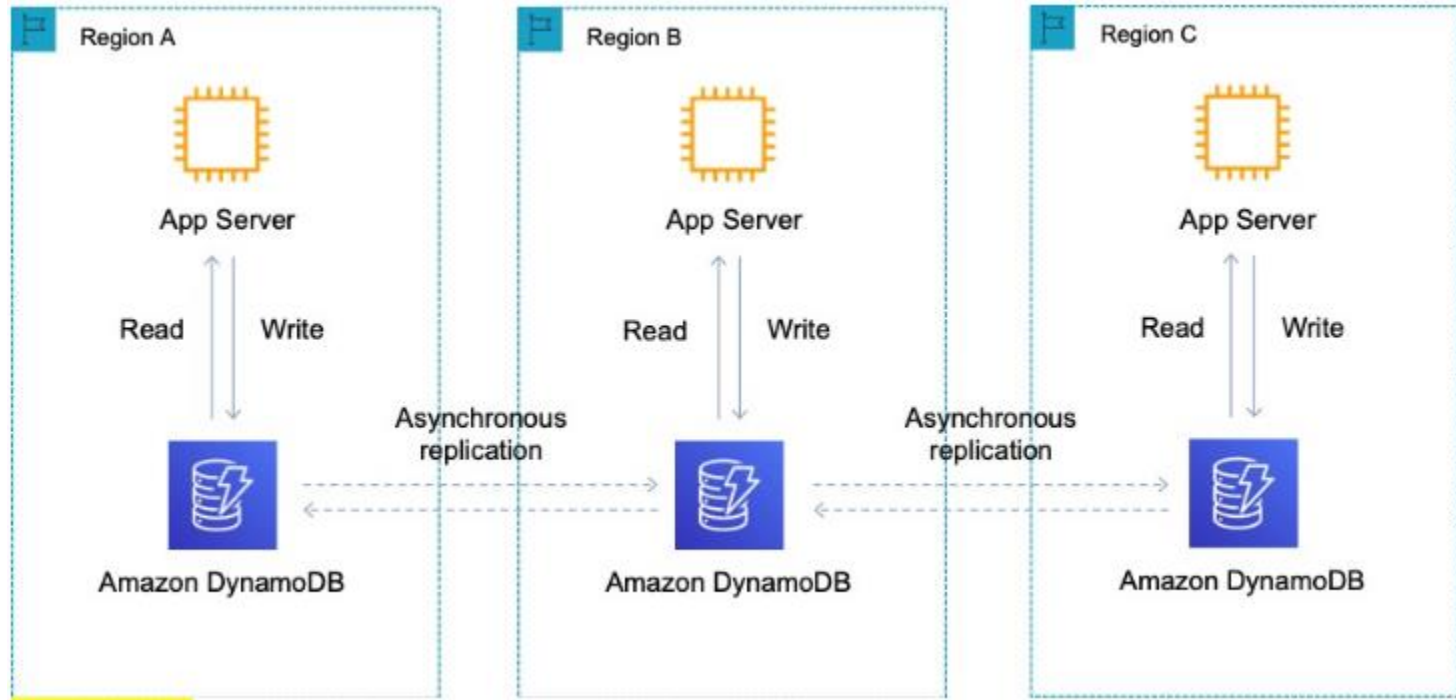


- Upto 10 GBPS
- VMDq
- TCP/IP
- Multiple ENI/instance
- Traffic can traverse across subnets
- VPC Networking, General purpose
- Default

- Upto 25 GBPS
- SR-IOV
- TCP/IP
- Single setting/per instance
- Traffic can traverses across subnets
- Low latency apps
- Optional on supported instance type

- Upto 100 GBPS
- OS-Bypass
- SRD
- One EFA per instance
- OS Bypass traffic is limited to single subnet and is not routable
- HPC and ML Apps
- Optional on supported instance type





* If your application requires **strongly consistent** reads, it must perform all of its strongly consistent reads and writes in the same Region. **DynamoDB does not support strongly consistent reads across Regions.**

Amazon EMR


Easily Run Spark, Hadoop, Hive, Presto, HBase, and more big data apps on AWS

Latest versions



Updated with latest open source frameworks within 30 days

Low cost




50–80% reduction in costs with EC2 Spot and Reserved Instances
Per-second billing for flexibility

Use S3 storage



Process data in S3 securely with high performance using the EMRFS connector

Easy



Fully managed no cluster setup, node provisioning, cluster tuning

S3 cross-region replication ^{New}

Automated, fast, and reliable asynchronous replication of data across AWS regions

Use cases

Compliance - store data hundreds of miles apart

Lower latency - distribute data to regional customers)

Security - create remote replicas managed by separate AWS accounts



- Only replicates new PUTs. Once S3 is configured, all new uploads into a source bucket will be replicated
- Entire bucket or prefix based
- 1:1 replication between any 2 regions
- Versioning required

Details on Cross-Region Replication

Versioning - Need to enable S3 versioning for the source and destination buckets.

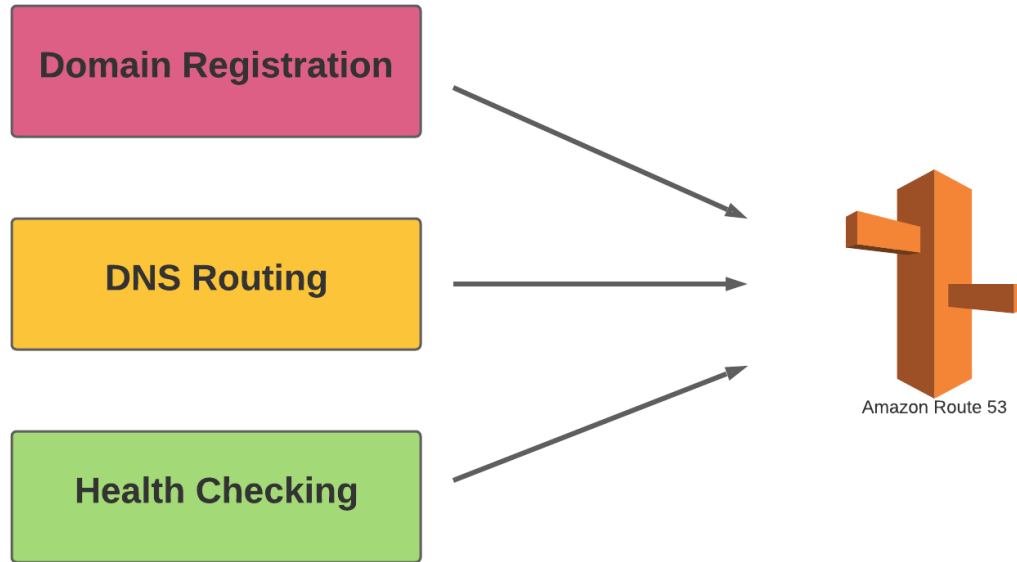
Lifecycle Rules - You can choose to use Lifecycle Rules on the destination bucket to manage older versions by deleting them or migrating them to Amazon Glacier.

Determining Replication Status - Use the HEAD operation on a source object to determine its replication status.

Region-to-Region - Replication always takes place between a pair of AWS regions. You cannot use this feature to replicate content to two buckets that are in the same region.

New Objects - Replicates new objects and changes to existing objects. Use S3 COPY to replicate existing objects

Amazon Route 53 is a highly available and scalable Domain Name System (DNS) web service.



Object lifecycle management

Transition actions—Define when objects transition to another storage class

Expiration actions—Define when objects expire. Amazon S3 deletes expired objects on your behalf.



Delete

Create trail



	Name ▲	Home region ▼	Multi-region trail ▼	Insights ▼	Organization trail ▼	S3 bucket ▼	Log file prefix ▼	CloudWatch Logs log group ▼	Status ▼
<input type="radio"/>	CloudTrail-event-log-files	US East (N. Virginia)	Yes	Disabled	No	aws-cloudtrail-logs-827784331229-b8c157f7			Logging

Trail name

Enter a display name for your trail.

3-128 characters. Only letters, numbers, periods, underscores, and dashes are allowed.

☐ Enable for all accounts in my organizationTo review accounts in your organization, open AWS Organizations. [See all accounts](#)Storage location [Info](#)☒ Create new S3 bucket
Create a bucket to store logs for this trail.☐ Use existing S3 bucket
Choose an existing bucket to store logs for this trail.

Trail log bucket and folder

Enter a new S3 bucket name and folder (prefix) to store your logs. Bucket names must be globally unique.

Logs will be stored in aws-cloudtrail-logs-827784331229-b8c157f7/AWSLogs/827784331229

Log file SSE-KMS encryption [Info](#)☒ Enabled

AWS KMS customer managed CMK

☒ New☐ Existing

AWS KMS alias

KMS key and S3 bucket must be in the same region.

▼ Additional settings

Log file validation [Info](#)☒ Enabled

1

Amazon S3 > aws-cloudtrail-logs-827784331229-b8c157f7 > AWSLogs/ > 827784331229/

827784331229/

3

Objects

Folder properties

Objects (2)

Objects are the fundamental entities stored in Amazon S3. For others to access your objects, you'll need to explicitly grant them permissions.




Delete


Actions ▼

Create folder

Upload

<input type="checkbox"/>	Name ▲	Type ▼	Last modified
<input type="checkbox"/>	CloudTrail-Digest/	Folder	-
<input type="checkbox"/>	CloudTrail/	Folder	-





- Each AZ consists of **one or more** physical data centers.

http://d1111111abcdef8.cloudfront.net/images/image.jpg?color=red&size=large

Select a delivery method for your content.

Web

Create a web distribution if you want to:

- Speed up distribution of static and dynamic content, for example, .html, .css, .php, and graphics files.
- Distribute media files using HTTP or HTTPS.
- Add, update, or delete objects, and submit data from web forms.
- Use live streaming to stream an event in real time.

You store your files in an origin - either an Amazon S3 bucket or a web server. After you create the distribution, you can add more origins to the distribution.

[Get Started](#)

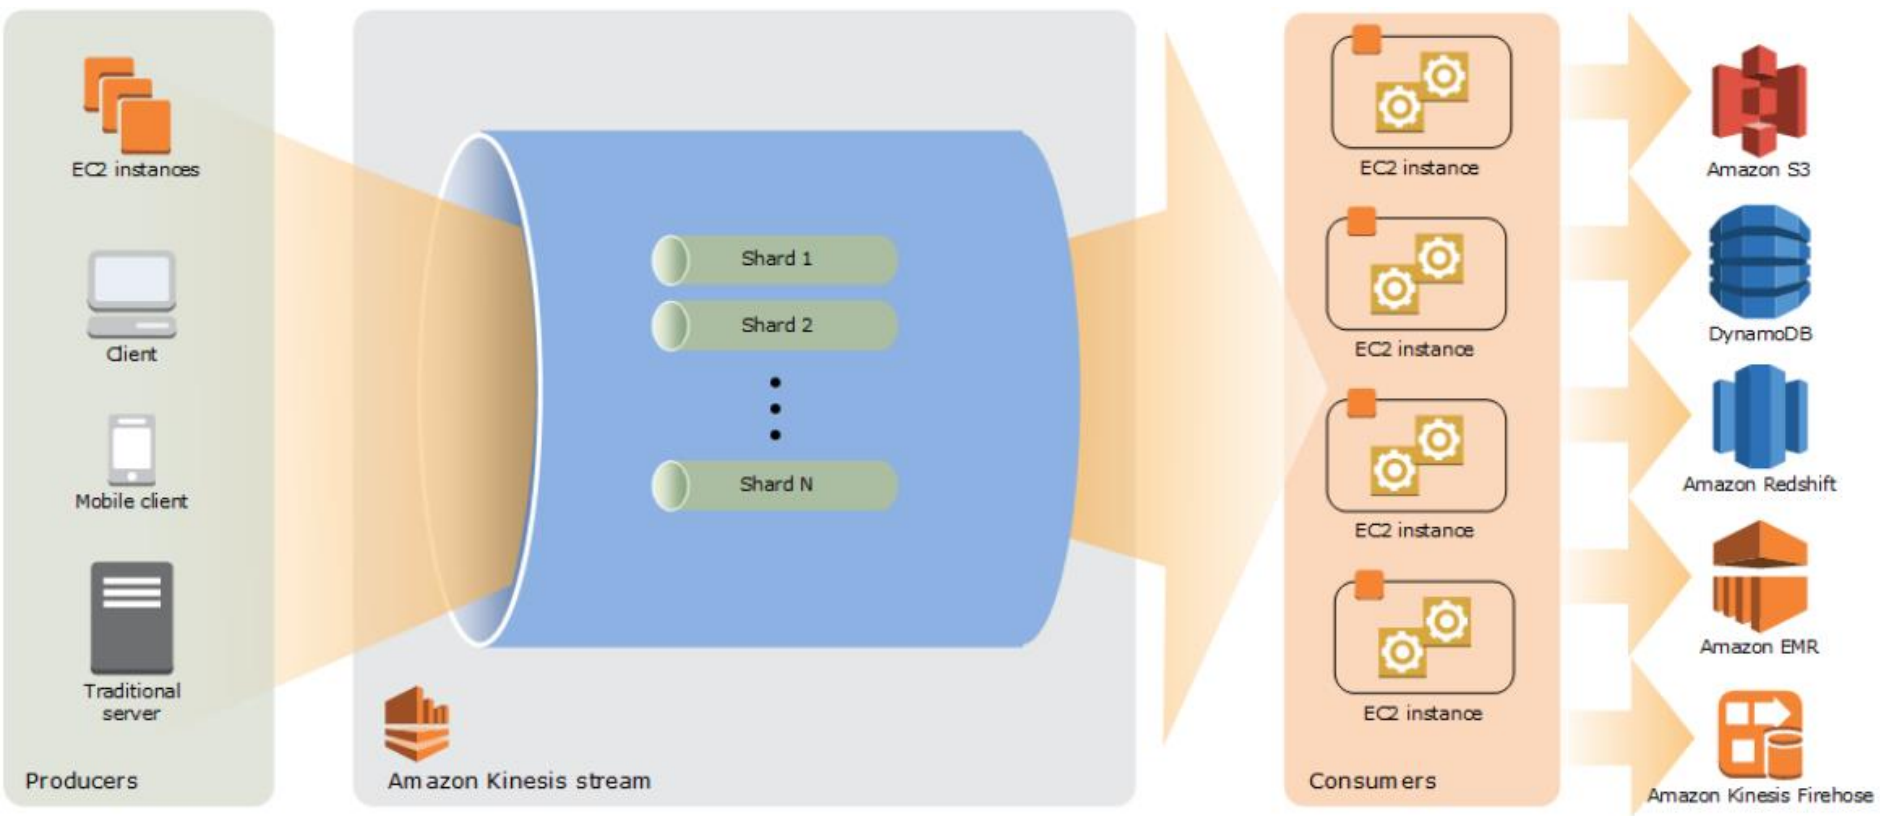
RTMP

CloudFront is discontinuing support for RTMP distributions on December 31, 2020. For more information, please [read the announcement](#).



Create an RTMP distribution to speed up distribution of your streaming media files using Adobe Flash Media Server's RTMP protocol. An RTMP distribution allows an end user to begin playing a media file before the file has finished downloading from a CloudFront edge location. Note the following:

- To create an RTMP distribution, you must store the media files in an Amazon S3 bucket.
- To use CloudFront live streaming, create a web distribution.

[Get Started](#)



spring.broadcast

Bucket overview			
Region	Amazon resource name (ARN)	Creation date	Access
US East (N. Virginia) us-east-1	 <code>arn:aws:s3:::spring.broadcast</code>	October 10, 2020, 16:21 (UTC+03:00)	 Public

Objects Properties Permissions Metrics Management Access points

Bucket Versioning

Versioning is a means of keeping multiple variants of an object in the same bucket. You can use versioning to preserve, retrieve, and restore every version of every object stored in your Amazon S3 bucket. With versioning, you can easily recover from both unintended user actions and application failures. [Learn more](#)

Edit

Bucket Versioning

Disabled

Multi-factor authentication (MFA) delete

An additional layer of security that requires multi-factor authentication for changing Bucket Versioning settings and permanently deleting object versions.

Disabled

Tags (0)

Track storage cost or other criteria by tagging your bucket. [Learn more](#)

Key

No tags

Amazon S3 > spring.broadcast > Edit Bucket Versioning

Edit Bucket Versioning

Bucket Versioning

Versioning is a means of keeping multiple variants of an object in the same bucket. You can use versioning to preserve, retrieve, and restore every version of every object stored in your Amazon S3 bucket. With versioning, you can easily recover from both unintended user actions and application failures. [Learn more](#)

Bucket Versioning

☐ Suspend

This suspends the creation of object versions for all operations but preserves any existing object versions.

☒ **Enable**

After enabling Bucket Versioning, you might need to update your lifecycle rules to manage previous versions of objects.

Multi-factor authentication (MFA) delete

An additional layer of security that requires multi-factor authentication for changing Bucket Versioning settings and permanently deleting object versions. To modify MFA delete settings, use the AWS CLI, AWS SDK, or the Amazon S3 REST API. [Learn more](#)

Disabled

Cancel

Save changes

STANDARD QUEUE

- **Unlimited Throughput:** Support a nearly unlimited number of transactions per second (TPS) per API action.
- **At-Least-Once Delivery:** A message is delivered at least once, but occasionally more than one copy of a message is delivered.
- **Best-Effort Ordering:** Occasionally, messages might be delivered in an order different from which they were sent.



FIFO QUEUE

- **High Throughput:** By default, FIFO queues support up to 300 messages per second
- **Exactly-Once Processing:** A message is delivered once and remains available until a consumer processes and deletes it. **Duplicates aren't introduced into the queue.**
- **First-In-First-Out Delivery:** The order in which messages are sent and received is strictly preserved





SQS - Short vs Long Polling

Polling is the method in which we retrieve messages from the queues.

Short polling (default) returns messages immediately, even if the message queue being polled is empty.

When you need a message **right away**, shorting polling is what you want to use.

Long polling waits until message **arrives in the queue**, or the **long poll timeout expires**.

Long polling makes it **inexpensive to retrieve messages** from your queue as soon as the messages are available.

Using long polling will reduce the cost because you can **reduce the number of empty receives**.

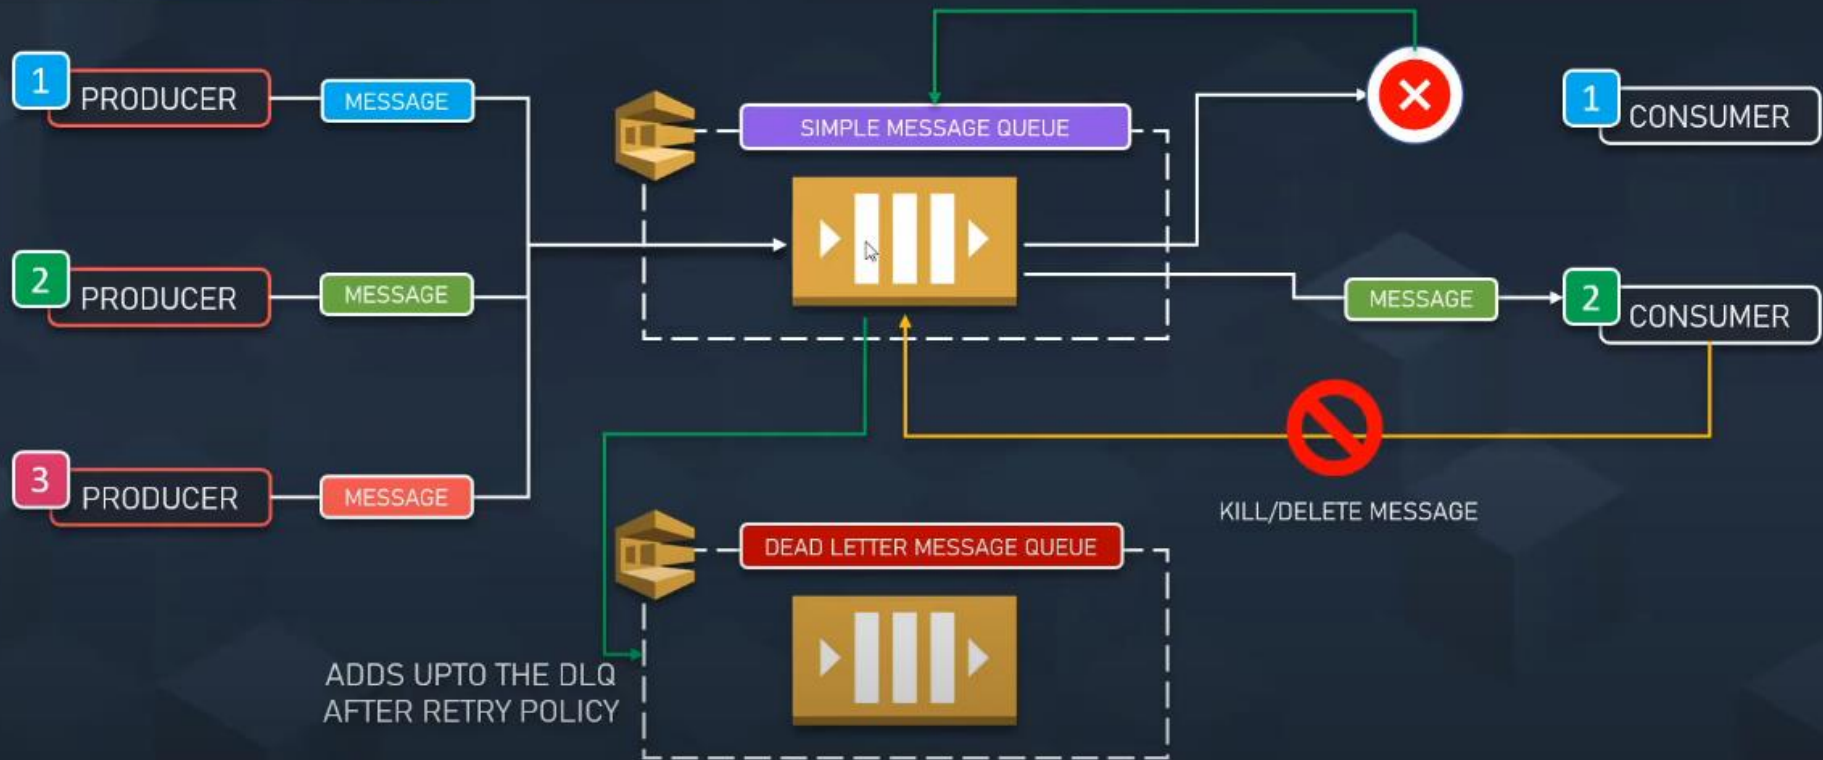
Most use-cases you want to use Long Polling

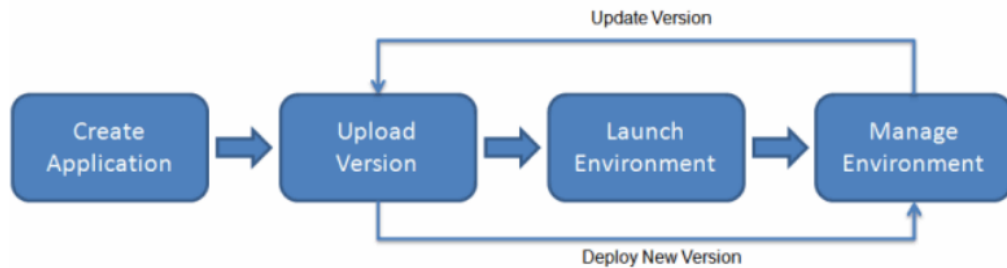
You can enable long polling when receiving a message by setting the wait time in seconds on the **ReceiveMessageRequest**



```
1 ReceiveMessageRequest receive_request = new ReceiveMessageRequest()
2   .withQueueUrl(queue_url)
3   .withWaitTimeSeconds(40);
4 sqs.receiveMessage(receive_request);
```

Dead Letter Queue





Java,
.NET,
PHP,
Node.js,
Python,
Ruby,
Go,
and Docker

Application information

Application name

Up to 100 Unicode characters, not including forward slash (/).

Application tags

Apply up to 50 tags. You can use tags to group and filter your resources. A tag is a key-value pair. The key must be unique within the resource and is case-sensitive. [Learn more](#)

Key Value

50 remaining

Platform

Platform

Platform branch

Platform version

Application code

☒ Sample application
Get started right away with sample code.

☐ Upload your code
Upload a source bundle from your computer or copy one from Amazon S3.

With Elastic Beanstalk, you can **quickly deploy** and manage applications in the AWS Cloud **without having to learn about the infrastructure** that runs those applications.

Use cases

Offload the SSL processing for web servers

Secure Sockets Layer (SSL) and Transport Layer Security (TLS) are used to confirm the identity of web servers and establish secure HTTPS connections over the Internet. You can use AWS CloudHSM to offload SSL/TLS processing for your web servers. Using CloudHSM for this processing reduces the burden on your web server and provides extra security by storing your web server's private key in CloudHSM.

Protect private keys for an issuing certificate authority (CA)

In a public key infrastructure (PKI), a certificate authority (CA) is a trusted entity that issues digital certificates. These digital certificates are used to identify a person or organization. You can use AWS CloudHSM to store your private keys and sign certificate requests so that you can securely act as an issuing CA to issue certificates for your organization.

Enable Transparent Data Encryption (TDE) for Oracle databases

You can use AWS CloudHSM to store the Transparent Data Encryption (TDE) master encryption key for your Oracle database servers that support TDE. Support for SQL Server is coming soon. With TDE, supported database servers can encrypt data before storing it on disk. Please note Amazon RDS for Oracle does not support TDE with CloudHSM; you should use AWS Key Management Service for this use case.

<https://docs.aws.amazon.com/cloudhsm/latest/userguide/backups.html>

Simple routing policy – basic routing policy defined using an A record to resolve to a single resource always without any specific rules.

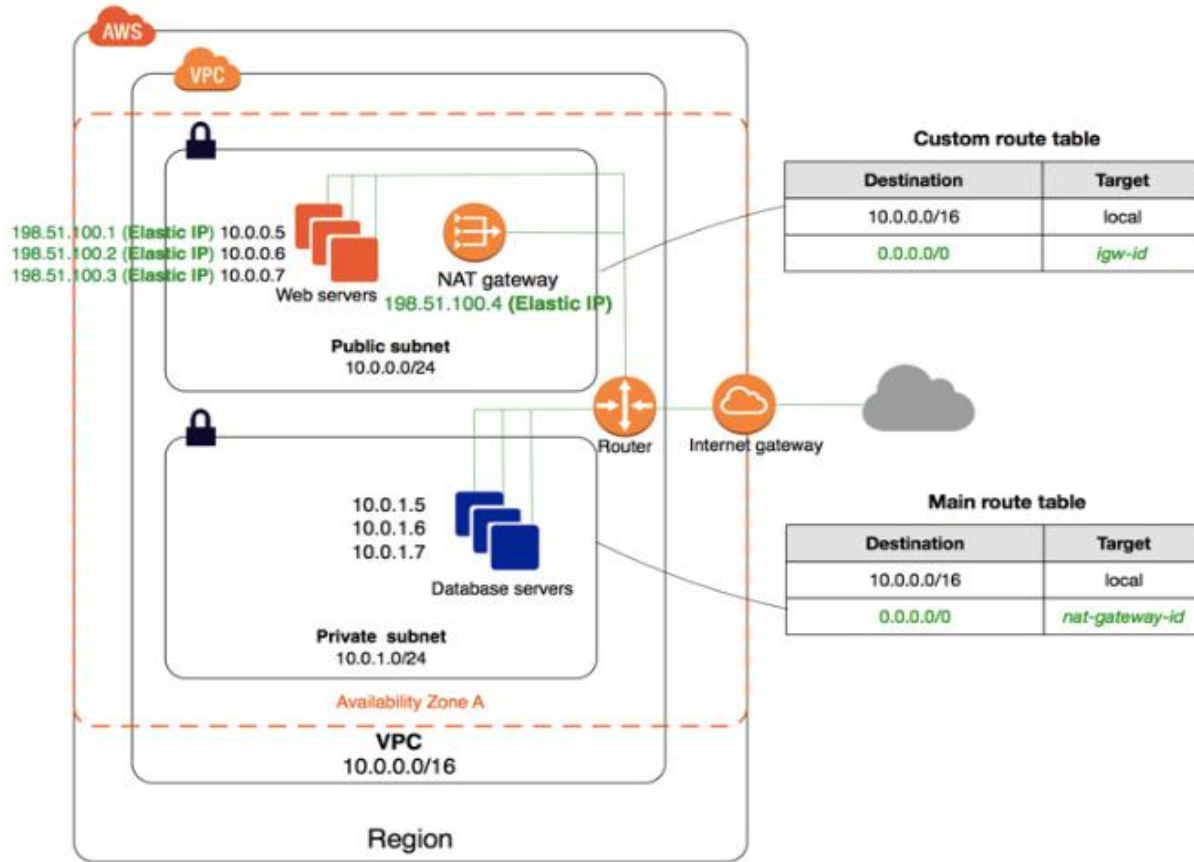
Multivalued answer routing policy – Use when you want Route 53 to respond to DNS queries with up to eight healthy records selected at **random**.

Latency routing policy – is used when there are multiple resources (multiple AWS Regions) for the **same functionality** and you want Route 53 to respond to DNS queries with answers that provide **the best latency**.

Weighted routing policy – is good for testing new versions of the software. Also, It is the ideal approach for **Blue-Green** deployments.

Solid State Drives (SSD)						Hard Disk Drives (HDD)	
Volume Type	EBS Provisioned IOPS SSD (io2 Block Express)	EBS Provisioned IOPS SSD (io2)	EBS Provisioned IOPS SSD (io1)	EBS General Purpose SSD (gp3) announced Dec 1, 2020	EBS General Purpose SSD (gp2)*	Throughput Optimized HDD (st1)	Cold HDD (sc1)
Short Description	Highest performance SSD volume designed for business-critical latency-sensitive transactional workloads	Highest performance and highest durability SSD volume designed for latency-sensitive transactional workloads	Highest performance SSD volume designed for latency-sensitive transactional workloads	Lowest cost SSD volume that balances price performance for a wide variety of transactional workloads	General Purpose SSD volume that balances price performance for a wide variety of transactional workloads	Low cost HDD volume designed for frequently accessed, throughput intensive workloads	Lowest cost HDD volume designed for less frequently accessed workloads
Durability	99.999%		99.8% - 99.9% durability			99.8% - 99.9% durability	
Use Cases	Largest, most I/O intensive, mission critical deployments of NoSQL and relational databases such as Oracle, SAP HANA, Microsoft SQL Server, and SAS Analytics	I/O-intensive NoSQL and relational databases	I/O-intensive NoSQL and relational databases	Virtual desktops, medium sized single instance databases such as Microsoft SQL Server and Oracle, latency sensitive interactive applications, boot volumes, and dev/test environments	Virtual desktops, medium sized single instance databases such as Microsoft SQL Server and Oracle, latency sensitive interactive applications, boot volumes, and dev/test environments	Big data, data warehouses, log processing	Colder data requiring fewer scans per day

Solid State Drives (SSD)						Hard Disk Drives (HDD)	
Volume Type	EBS Provisioned IOPS SSD (io2 Block Express)	EBS Provisioned IOPS SSD (io2)	EBS Provisioned IOPS SSD (io1)	EBS General Purpose SSD (gp3)	EBS General Purpose SSD (gp2)*	Throughput Optimized HDD (st1)	Cold HDD (sc1)
API Name	io2	io2	io1	gp3	gp2	st1	sc1
Volume Size	4 GB – 64 TB			1 GB - 16 TB		125 GB - 16 TB	
Max IOPS**/Volume	256,000	64,000	64,000	16,000	16,000	500	250
Max Throughput***/Volume	4,000 MB/s	1,000 MB/s	1,000 MB/s	1,000 MB/s	250 MB/s	500 MB/s	250 MB/s
Max IOPS/Instance	260,000	160,000**	260,000	260,000	260,000	260,000	260,000
Max Throughput/Instance	7,500 MB/s	4,750 MB/s**	7,500 MB/s	7,500 MB/s	7,500 MB/s	7,500 MB/s	7,500 MB/s
Latency	sub-millisecond	single digit millisecond					
Price	\$0.125/GB-month \$0.065/provisioned IOPS-month up to 32,000 IOPS \$0.046/provisioned IOPS-month from 32,001 to 64,000 \$0.032/provisioned IOPS-month for greater than 64,000 IOPS		\$0.125/GB-month \$0.065/provisioned IOPS-month	\$0.08/GB-month 3,000 IOPS free and \$0.005/provisioned IOPS-month over 3,000; 125 MB/s free and \$0.04/provisioned MB/s-month over 125	\$0.10/GB-month	\$0.045/GB-month	\$0.015/GB-month
Dominant Performance Attribute	IOPS, throughput, latency, capacity, and volume durability	IOPS and volume durability	IOPS	IOPS	IOPS	MB/s	MB/s



partition key

Music

```
{  
  "Artist": "No One You Know",  
  "SongTitle": "My Dog Spot",  
  "AlbumTitle": "Hey Now",  
  "Price": 1.98,  
  "Genre": "Country",  
  "CriticRating": 8.4  
}
```

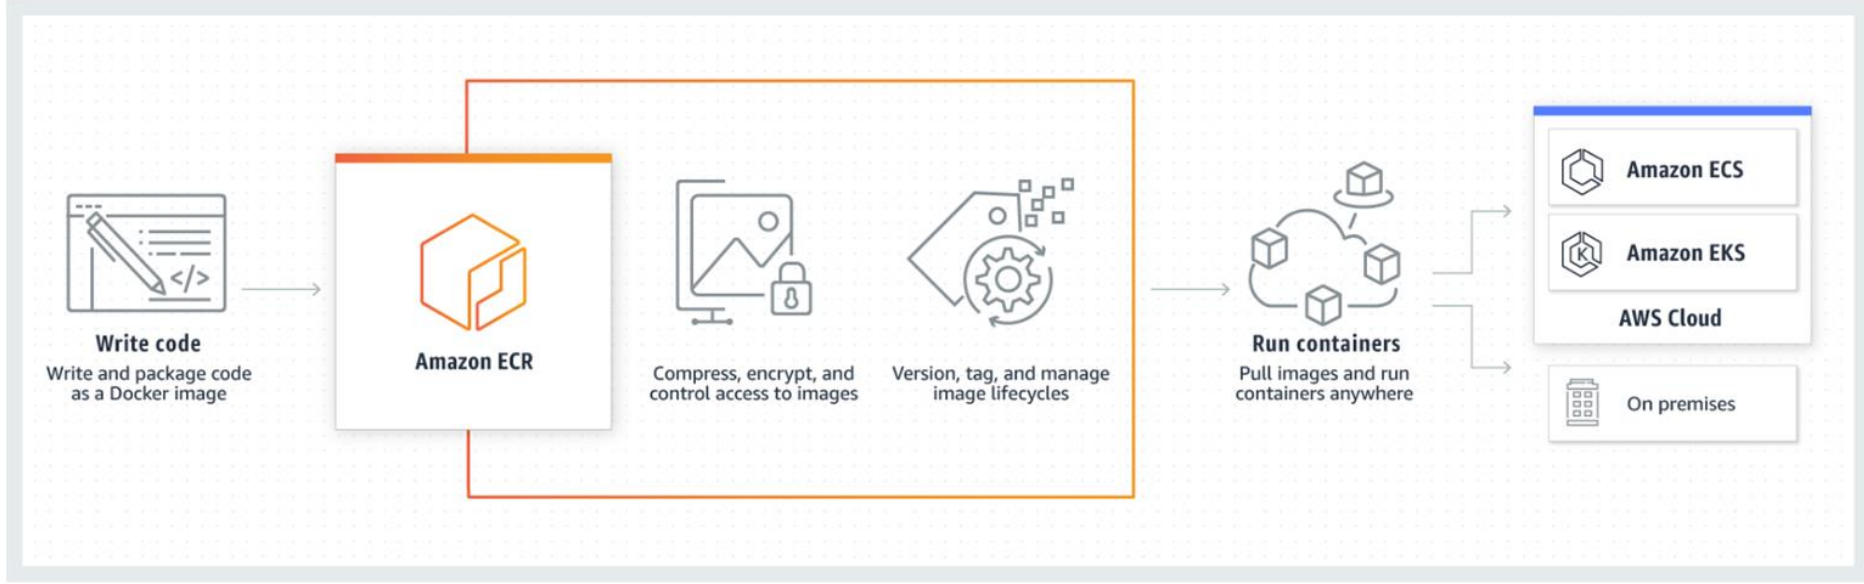
sort key

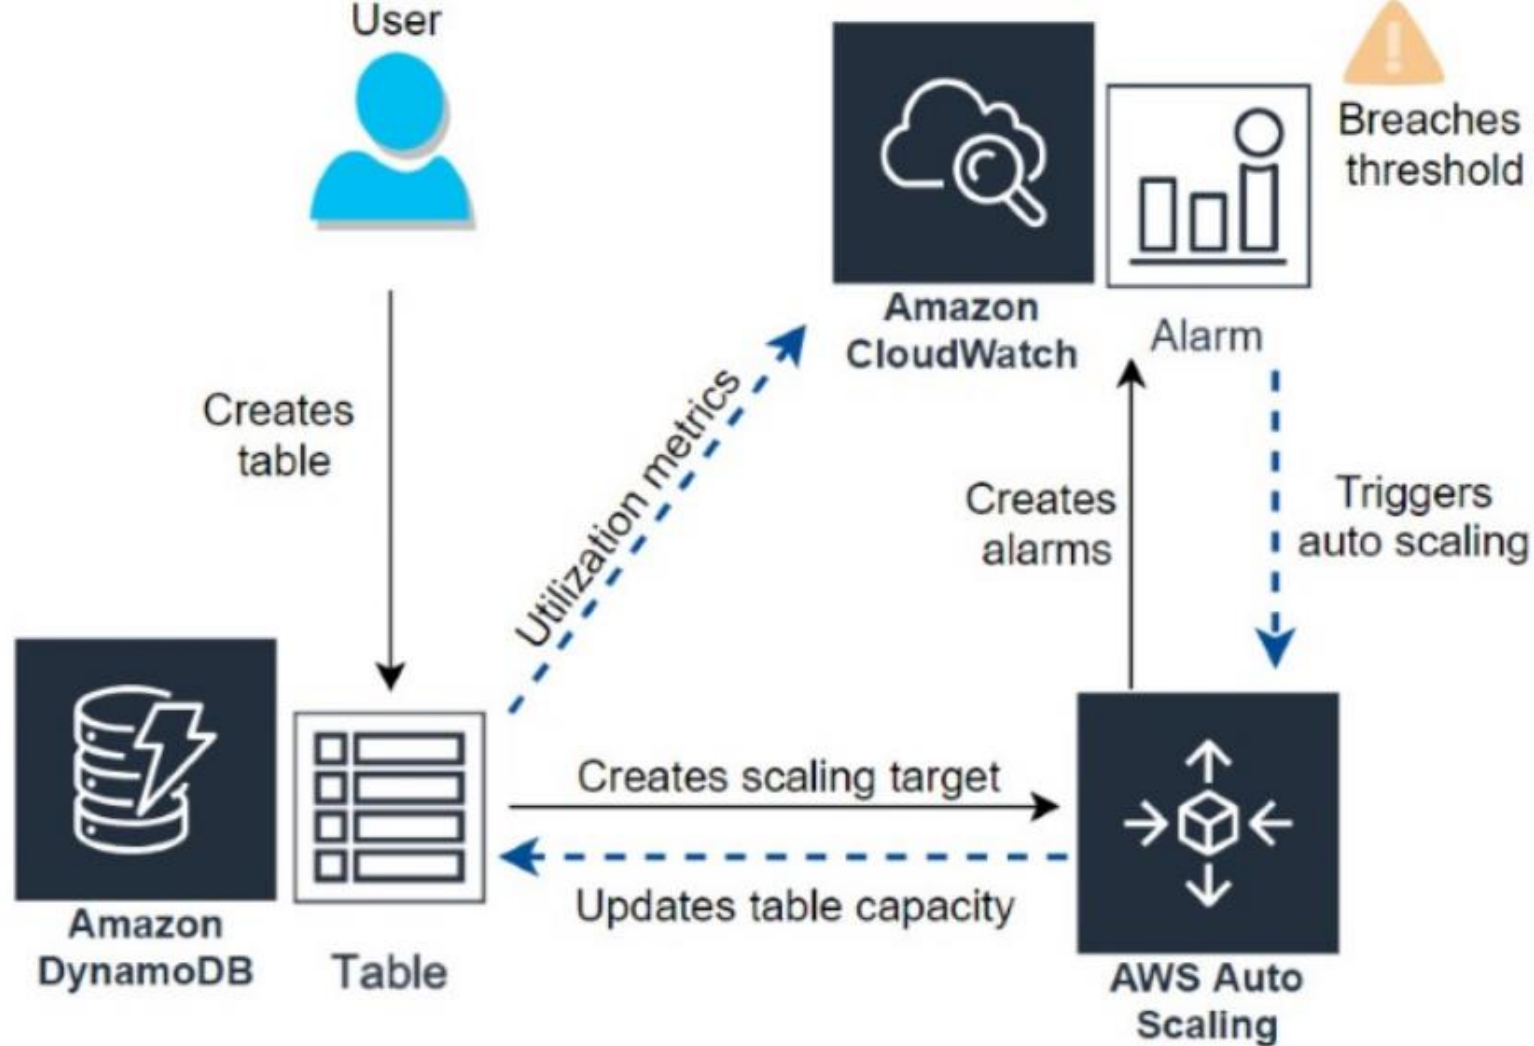
```
{  
  "Artist": "No One You Know",  
  "SongTitle": "Somewhere Down The Road",  
  "AlbumTitle": "Somewhat Famous",  
  "Genre": "Country",  
  "CriticRating": 8.4,  
  "Year": 1984  
}
```

GenreAlbumTitle

```
{  
  "Genre": "Country",  
  "AlbumTitle": "Hey Now",  
  "Artist": "No One You Know",  
  "SongTitle": "My Dog Spot"  
}
```

```
{  
  "Genre": "Country",  
  "AlbumTitle": "Somewhat Famous",  
  "Artist": "No One You Know",  
  "SongTitle": "Somewhere Down The Road"  
}
```





Create Distribution

Origin Settings

Origin Domain Name


Origin Path

Enable Origin Shield ☐ Yes
☒ No

Origin ID

Restrict Bucket Access ☒ Yes
☐ No

Origin Access Identity ☒ Create a New Identity
☐ Use an Existing Identity



signed URLs or
signed cookies
(https only)

1

i


i

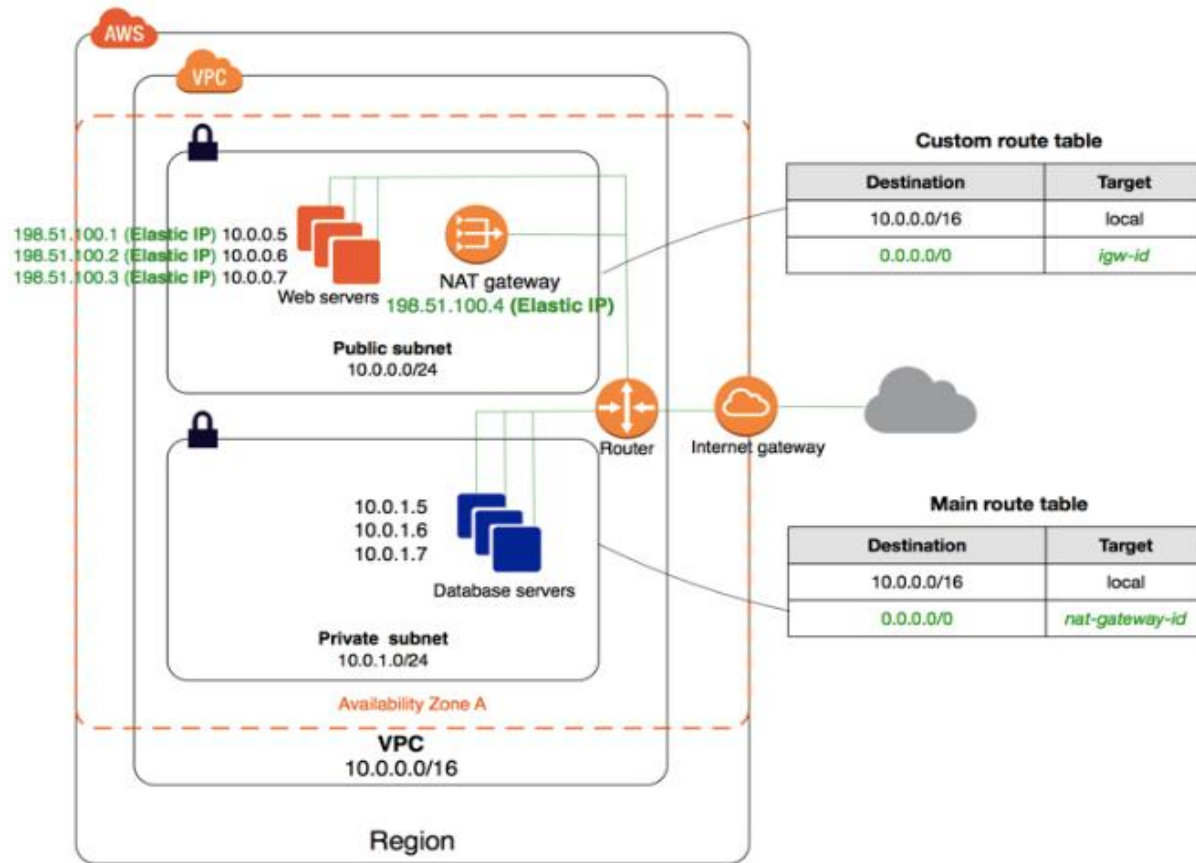
Enter a description for the origin. This value lets you distinguish multiple origins in the same distribution from one another. The description for each origin must be unique within the distribution.

i

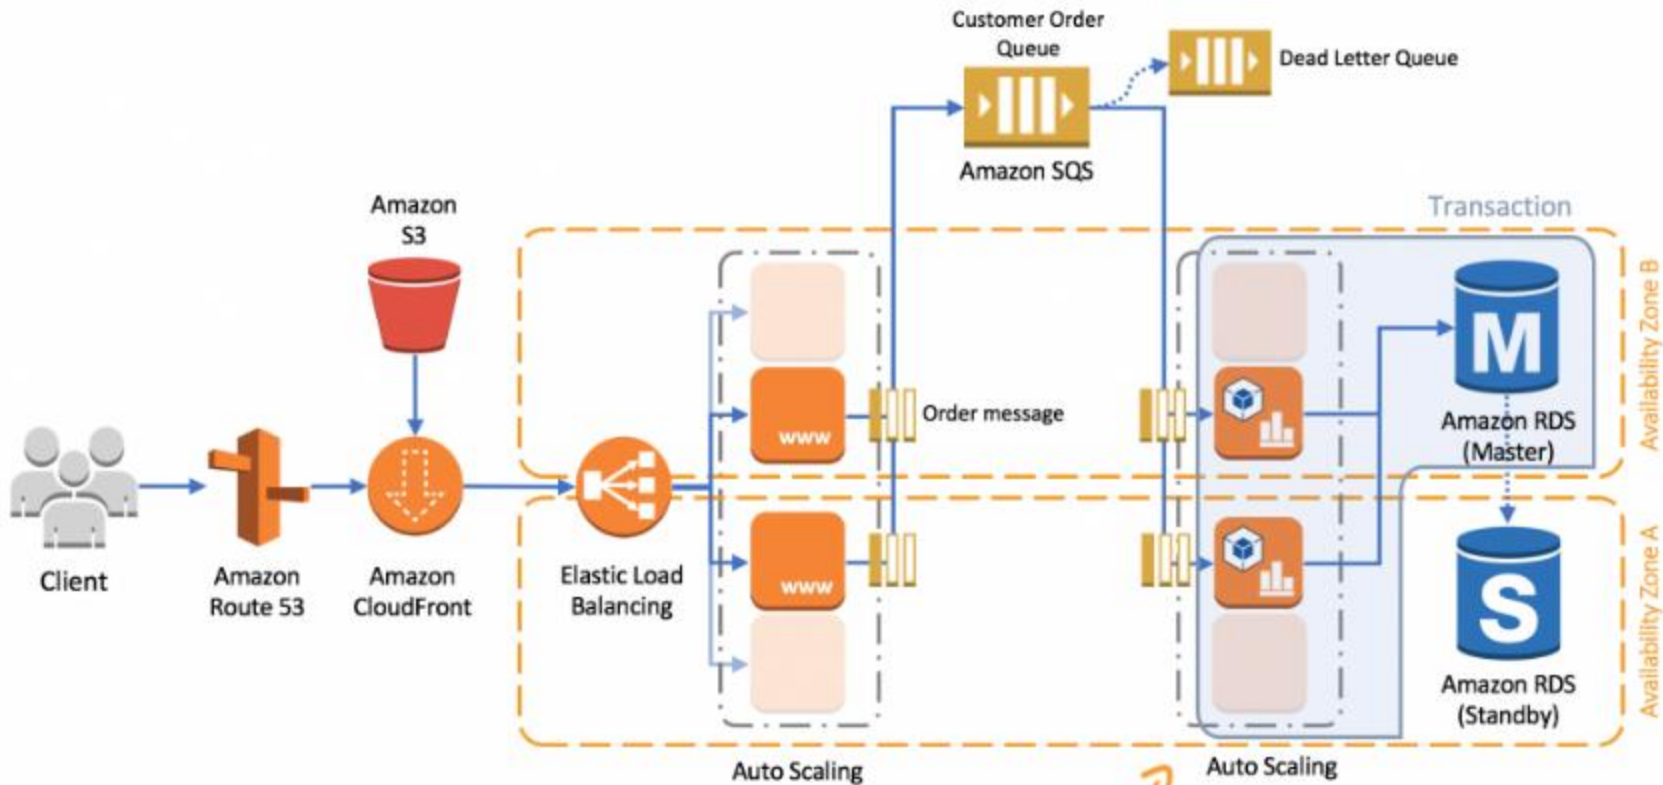
If you want to require that users always access your Amazon S3 content using CloudFront URLs, not Amazon S3 URLs, click Yes. This is useful when you are using signed URLs or signed cookies to restrict access to your content. In the Help, see "Serving Private Content through CloudFront".

i










New processing nodes scaling independently, using

- ApproximateNumberOfMessagesVisible
- ApproximateAgeOfOldestMessage

us-east-1 region

a	b	c	d	e	f	Total	
3	3	3	3			12	
16						16	
2	2	2	2	2		10	
4	4	4				12	

SNS is a distributed **publish-subscribe** system. Messages are pushed to subscribers when they are sent by publishers to SNS. AWS SNS is able to push notifications to the related **SQS endpoints**.



Objects

Properties

Permissions

Metrics

Management


Access points

Bucket overview

Region

US East (N. Virginia) us-east-1


Amazon resource name (ARN)

arn:aws:s3:::deneme

Creation date

September 19, 2020, 22:21:02 (UTC+03:00)

Bucket Versioning


Versioning is a means of keeping multiple variants of an object in the same bucket. You can use versioning to preserve, retrieve, and restore every version of every object stored in your Amazon S3 bucket. With versioning, you can easily recover from both unintended user actions and application failures. [Learn more](#) 

Edit

Bucket Versioning

Suspended

Multi-factor authentication (MFA) delete

An additional layer of security that requires multi-factor authentication for changing Bucket Versioning settings and permanently deleting object versions. To modify MFA delete settings, use the AWS CLI, AWS SDK, or the Amazon S3 REST API. [Learn more](#) 

Disabled

Object Lock					<div>Edit</div>
Store objects using a write-once-read-many (WORM) model to help you prevent objects from being deleted or overwritten for a fixed amount of time or indefinitely. Learn more					
Object Lock					
Disabled					
<div><div></div> Amazon S3 currently does not support enabling Object Lock after a bucket has been created. To enable Object Lock for this bucket, contact Customer Support</div>					

Comparing the Amazon S3 storage classes

The following table compares the storage classes.

Storage class	Designed for	Durability (designed for)	Availability (designed for)	Availability Zones	Min storage duration
S3 Standard	Frequently accessed data	99.999999999%	99.99%	>= 3	None
S3 Standard-IA	Long-lived, infrequently accessed data	99.999999999%	99.9%	>= 3	30 days
S3 Intelligent-Tiering	Long-lived data with changing or unknown access patterns	99.999999999%	99.9%	>= 3	30 days
S3 One Zone-IA	Long-lived, infrequently accessed, <u>non-critical data</u>	99.999999999%	99.5%	1	30 days
S3 Glacier	Long-term data archiving with retrieval times ranging from minutes to hours	99.999999999%	99.99% (after you restore objects)	>= 3	90 days
S3 Glacier Deep Archive	Archiving rarely accessed data with a default retrieval time of 12 hours	99.999999999%	99.99% (after you restore objects)	>= 3	180 days
RRS (Not recommended)	Frequently accessed, non-critical data	99.99%	99.99%	>= 3	None